

FERPA & AI Tools in Teaching

What You Can — and Can't — Feed an AI

Emily Beam

University of Vermont · Thinking with Agents

Teaching Webinar · June 17, 2026

The 30-Second Version

FERPA applies the moment you put identifiable student work, grades, or education records into an AI tool. There are exactly **two** clean, legal ways to use AI with student data.

✓ Path 1: Approved tool

Use an **institution-vetted** AI tool that qualifies as a “school official” (under contract / institutional control).

✓ Path 2: De-identify

Properly de-identify the data first, so it's no longer protected information.

⊘ The one-line don't

Never paste identifiable student work or grades into your **personal/consumer** AI account (consumer ChatGPT, Claude, or Copilot). That's an unconsented disclosure of education records.

Not legal advice — accurate to 34 CFR Part 99 & ED guidance as of June 2026; defer to your institution's policy.

Part 1

The Legal Foundation

What FERPA Actually Requires

What's Protected: "Education Records"

Records that are:

1. **directly related to a student, AND**
2. **maintained by the institution** — or by a **party acting for it**.

Records an AI vendor holds *on your institution's behalf* are **themselves** education records. The data stays under FERPA after you hand it over.

Grades, transcripts, class lists, and student work are the canonical examples.

34 CFR 99.3; ED "Responsibilities of Third-Party Service Providers under FERPA" (2015).

Useful nuance:

A *student* pasting *their own* work into a consumer tool does **not** make the vendor's records "education records."

The obligation runs to **you and the institution** — not the student.

What's Identifiable: PII Is Broader Than Names



FERPA's PII includes...

- ▶ Direct identifiers (name, SSN, student ID)
- ▶ **Indirect** identifiers (DOB, place of birth)
- ▶ Interaction **metadata**
- ▶ Even **aggregate** data

...anything that, alone or combined, lets “a reasonable person in the school community identify the student with reasonable certainty.”



The mosaic effect

In a **small economics seminar**, an “anonymized” essay can still be re-identifiable from:

- ▶ topic
- ▶ writing style
- ▶ other facts you already know

De-identification means addressing **combinations** of identifiers — not just deleting names.

The Default Rule: Consent

A school **may not disclose PII** from education records **without prior signed, dated written consent** — *unless* an enumerated exception applies.

“Disclosure” explicitly covers electronic transmission “by any means . . . to any party.” **Uploading to an external AI tool is a disclosure.**

34 CFR 99.30(a), 99.3; ED 2014 online-services guidance.

In higher ed, the consent is the *student's* — not the parent's.

FERPA rights transfer to the student when they turn 18 **or enter a post-secondary institution at any age.**

Being a “disclosure” $\neq$ automatically a violation — it means an exception or de-identification must apply.

Exception A: The “School Official” Pathway (for vendors)

A vendor may receive PII **without consent** only if it:

1. performs a service the school would otherwise use **employees** for;
2. is a designated school official with a **legitimate educational interest** (annual FERPA notice);
3. is under the institution’s **direct control**; and
4. uses records **only for authorized purposes — no re-disclosure**.

34 CFR 99.31(a)(1)(i); 99.33(a); ED Vendor FAQ + 2014 guidance.



Bottom line

A **consumer/personal** AI account is **not** under institutional control. It **can-not** qualify.

Only an institutionally-contracted, vetted tool can.

This same exception is *why you* can lawfully see your own students’ records — instructors are “school officials with legitimate educational interests.”

Exception B: De-Identification (no vendor contract needed)

After removing all PII, an institution may use/release records without consent — *provided* it makes a **reasonable determination** that identity is not personally identifiable, accounting for single *or multiple* releases and **other reasonably available information**.

Properly de-identified data is **outside FERPA entirely**, so it may be used in general AI tools.

34 CFR 99.31(b)(1); ED 2014 guidance.



Say this out loud

“Properly de-identified” is a **risk-based standard**, not a name-stripping checklist.

In small classes the bar is **genuinely hard to clear**. Treat it as a judgment call.

What the Dept of Education Has — and Hasn't — Said



No AI-specific rules

No AI-specific FERPA regulations exist (as of June 2026). The existing **technology-neutral** guidance governs AI vendors directly.

The framework (school-official exception, direct control, data-minimization) is tech-neutral — AI tools are just a subset of “online educational services.”

The operative ED documents:

- ▶ Third-Party Service Providers under FERPA (2015)
- ▶ Online Educational Services: Requirements & Best Practices (2014)
- ▶ **Model Terms of Service checklist (2016)** — ready-made for vetting a tool's ToS

studentprivacy.ed.gov

Part 2

Practical Do's & Don'ts
For Economics Faculty

Do's & Don'ts

✗ DON'T

- ▶ Paste **identifiable** essays, exams, or grades into a **consumer** AI account
- ▶ Assume **deleting the name** makes an essay safe
- ▶ Upload a **gradebook** / **roster** to “summarize”
- ▶ Build a chatbot **trained on** identifiable student work without approval + a DPA

✓ DO

- ▶ Use your institution's **vetted, contracted** AI tool for identifiable data
- ▶ **De-identify first** — for re-ID risk, not just names
- ▶ Run a tool's ToS through ED's **2016 Model ToS checklist**
- ▶ Check your **institution's own AI policy** (often stricter)
- ▶ Ask your **registrar** / **privacy office** when unsure

Concrete Classroom Scenarios

Scenario	Verdict	Why
Named essay into consumer ChatGPT for feedback	✗	Unconsented disclosure
Same essay, fully de-identified , low re-ID risk, general tool	⚠	Defensible — your judgment
Same essay, in the university's contracted tool	✓	Vendor is a school official
Upload the class gradebook to summarize performance	✗	Grades are PII
Course chatbot trained on student submissions	✗	Needs approval + DPA
AI to draft a generic rubric or lecture (no student data)	✓	No education records involved

The dividing line is almost always: **identifiable student data + a tool that isn't institutionally controlled.**

The Question Everyone Asks: “Does It Train on My Data?”



The FERPA answer

Under the school-official exception, a vendor may use PII **only for the contracted purpose** and may not reuse it for its own ends.

So training a vendor's models on **identifiable** student records is an **unauthorized reuse** — unless you authorize it.

De-identified data is outside FERPA — training on it is fine.

ED 2014/2016 guidance; CRS R46799. A sound inference from the use-limitation rule — ED has no AI-specific ruling.



The vendor reality

- ▶ OpenAI, MS 365 Copilot, and Claude (**commercial/edu** tiers) all state they don't train on your prompts *by default*.
- ▶ These are **vendor statements**, not FERPA rulings — and consumer tiers are excluded.
- ▶ Only **Microsoft** publishes an explicit FERPA “school-official” commitment.

Full three-way comparison on the final (reference) slide.

So How Worried Should You Be? The Enforcement Reality



How FERPA is enforced

- ▶ **No private right of action** — a student cannot sue you under FERPA (*Gonzaga v. Doe, 2002*).
- ▶ ED's ultimate stick is **withdrawing federal funds** — never actually used.
- ▶ Real ed-tech enforcement has run through **COPPA/FTC**, not FERPA (Edmodo, \$6M, 2023).

So the real exposure is institutional and reputational, not a lawsuit.

That's exactly why your university **vets and contracts** tools centrally — and why the safe move is to **use the approved one**.

ED's only recent vendor action (the May 2026 Canvas/Instructure letter) was about a **data breach** — not AI. There is still **no AI-specific FERPA guidance**.

Takeaways: A Faculty Checklist

Before you put student data in any AI tool

1. Is it **identifiable** student data? (Remember: PII is broader than names.)
2. If yes — is the tool **institutionally vetted** / **contracted**? If not, **stop**.
3. Or can you **properly de-identify** it (risk-based, hard in small classes)?
4. When in doubt, run the ToS through ED's **Model ToS checklist** and ask your **privacy office**.

The one thing to remember: approved tool *or* de-identify — never identifiable student data in a personal AI account.

Reference: The Three Enterprise Tools, Compared

	ChatGPT Edu / Enterprise	MS 365 Copilot	Claude Edu / Enterprise
Trains on your data by default?	No	No	No
Consumer tier trains by default?	Yes (opt-out)	N/A	Yes since Aug 2025
Explicit FERPA “school official” language?	No	Yes (99.33(a))	No (processor + DPA)
Data Processing Addendum?	Yes	Yes (OST DPA)	Yes (auto-incl.)
Retention	Configurable	Configurable	Until deleted (API: 30d)
Zero-data-retention option?	Via API	—	API only, not chat

All entries are **vendor representations**, not adjudicated FERPA compliance — terms change fast. **Microsoft** alone publishes an explicit FERPA school-official commitment; the others rely on the contract/DPA + institutional vetting. *Sources: openai.com, learn.microsoft.com, privacy.claude.com — as of June 2026.*